

“I need to sort out my wardrobe”

Deep Learning Demystified



Amy Boyd - Cloud Developer Advocate AI/ML
@AmyKateNicho



“Deep Learning is not
as scary as I thought it was”

“I now feel I can hold a
good conversation with a
Data Scientist”

“I want to learn more and
give it a go myself”

Sound



Image



Video




Never Stop
DREAMING




Text

Traditional ML: Neural Network



Deep Learning: Deep Neural Network



Sound



Image



Video



Never Stop
DREAMING



Text

Computers are so Literal!

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1
-1	-1	-1	1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1


Computers are so Literal!



Computers are so Literal!



CNN: Convolutional Neural Network



 Conv2D

 MaxPooling

 ReLu


 Dense



Zalando

Fashion-MNIST Dataset

<https://github.com/zalandoresearch/fashion-mnist>



```
In [9]: #compile - how to measure loss
model.compile(loss=losses.categorical_crossentropy, optimizer=optimizers.Adam(), metrics=['accuracy'])

#train the model and return loss and accuracy for each epoch - history dictionary
start = time.time()
hist = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
end = time.time()

#evaluate the model on the test data
score = model.evaluate(x_test, y_test, verbose=0)
print('Test Loss: ', score[0])
print('Test Accuracy: ', score[1])
print('Time to run: ', (end-start))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/5

60000/60000 [=====] - 113s 2ms/step - loss: 0.5295 - acc: 0.8104 - val_loss: 0.3782 - val_acc: 0.8639

Epoch 2/5

60000/60000 [=====] - 104s 2ms/step - loss: 0.3389 - acc: 0.8779 - val_loss: 0.3500 - val_acc: 0.8784

Epoch 3/5

60000/60000 [=====] - 108s 2ms/step - loss: 0.2947 - acc: 0.8924 - val_loss: 0.3135 - val_acc: 0.8861

Epoch 4/5

60000/60000 [=====] - 127s 2ms/step - loss: 0.2659 - acc: 0.9016 - val_loss: 0.2789 - val_acc: 0.8977

Epoch 5/5

60000/60000 [=====] - 117s 2ms/step - loss: 0.2457 - acc: 0.9094 - val_loss: 0.2664 - val_acc: 0.9030

Test Loss: 0.26641942536830904

Test Accuracy: 0.903

Time to run: 569.7223751544952

```

In [7]: #compile - how to measure loss
model.compile(loss=losses.categorical_crossentropy, optimizer=optimizers.Adam(), metrics=['accuracy'])

#train the model and return loss and accuracy for each epoch - history dictionary
start = time.time()
hist = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
end = time.time()

#evaluate the model on the test data
score = model.evaluate(x_test, y_test, verbose=0)
print('Test Loss: ', score[0])
print('Test Accuracy: ', score[1])
print('Time to run: ', (end-start))

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 31s 509us/step - loss: 0.5248 - acc: 0.8110 - val_loss: 0.3996 - val_acc: 0.8558
Epoch 2/10
60000/60000 [=====] - 28s 467us/step - loss: 0.3372 - acc: 0.8788 - val_loss: 0.3200 - val_acc: 0.8870
Epoch 3/10
60000/60000 [=====] - 29s 481us/step - loss: 0.2904 - acc: 0.8944 - val_loss: 0.3582 - val_acc: 0.8709
Epoch 4/10
60000/60000 [=====] - 27s 456us/step - loss: 0.2614 - acc: 0.9040 - val_loss: 0.2764 - val_acc: 0.9014
Epoch 5/10
60000/60000 [=====] - 27s 452us/step - loss: 0.2345 - acc: 0.9147 - val_loss: 0.2665 - val_acc: 0.9038
Epoch 6/10
60000/60000 [=====] - 34s 559us/step - loss: 0.2183 - acc: 0.9191 - val_loss: 0.2569 - val_acc: 0.9082
Epoch 7/10
60000/60000 [=====] - 34s 562us/step - loss: 0.1976 - acc: 0.9279 - val_loss: 0.2458 - val_acc: 0.9102
Epoch 8/10
60000/60000 [=====] - 27s 457us/step - loss: 0.1813 - acc: 0.9334 - val_loss: 0.2641 - val_acc: 0.9064
Epoch 9/10
60000/60000 [=====] - 28s 469us/step - loss: 0.1690 - acc: 0.9376 - val_loss: 0.2578 - val_acc: 0.9063
Epoch 10/10
60000/60000 [=====] - 27s 448us/step - loss: 0.1542 - acc: 0.9432 - val_loss: 0.2465 - val_acc: 0.9134
Test Loss: 0.2464900684028864
Test Accuracy: 0.9134
Time to run: 292.27598786354065

```

*** free compute*

CPU

Avg = ~114 secs

*Time to run per
epoch. 5 epochs run*

*** NC6 series*

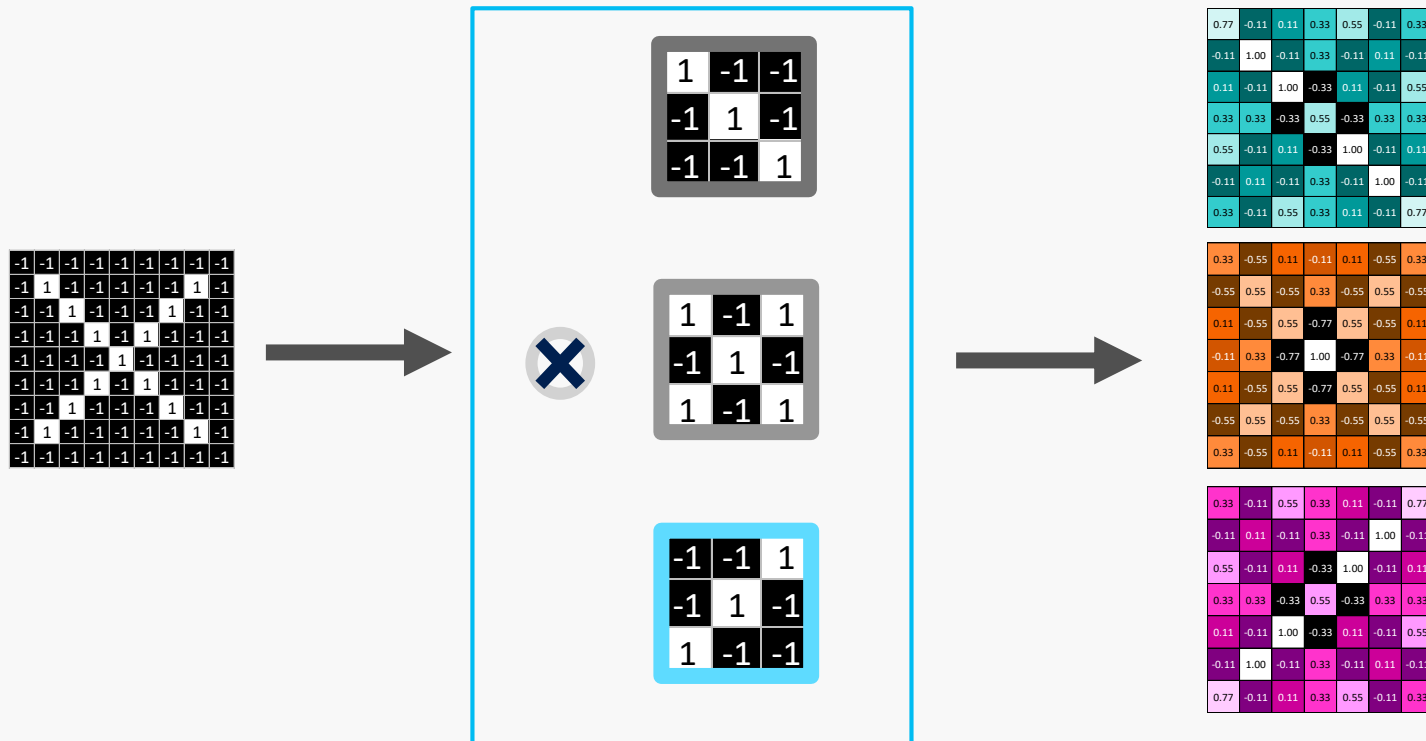
GPU

Avg = ~29 secs

*Time to run per
epoch. 10 epochs run*







Convolution layer

One image becomes a stack of filtered images



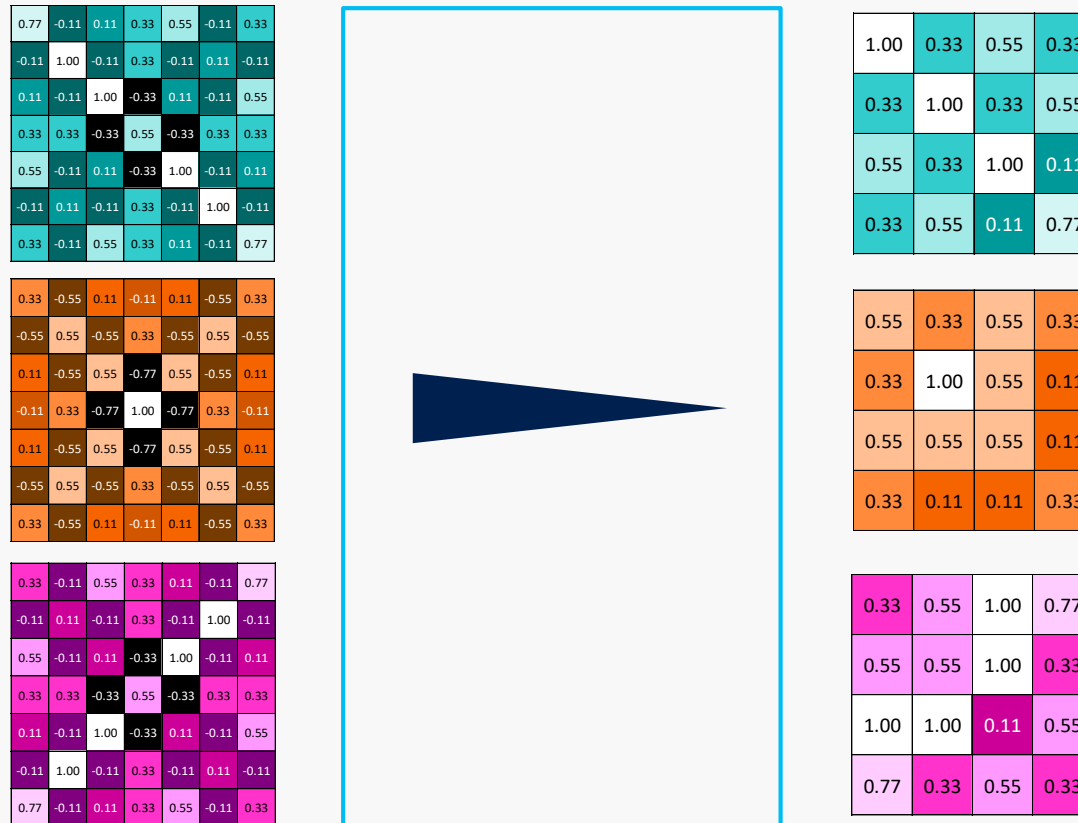
ReLU layer

A stack of images becomes a stack of images with no negative values.



Pooling layer

A stack of images becomes a stack of smaller images.



Deep Learning: Key Takeaways

- ⊗ Deep Learning is not as scary as I thought it was
- ⊗ I now feel I can hold a good conversation with a Data Scientist
- ⊗ I want to learn more and give it a go myself

Deep Learning: Key Takeaways

- ⊗ **Azure Notebooks:** <https://docs.microsoft.com/en-us/azure/notebooks/>
- ⊗ **Deep learning with Keras Course:** <https://app.pluralsight.com/library/courses/keras-deep-learning/table-of-contents>
- ⊗ **Keras API Documentation:** <https://keras.io/>

“I need to sort out my wardrobe”

Deep Learning Demystified



Amy Boyd - Cloud Developer Advocate AI/ML
@AmyKateNicho

